

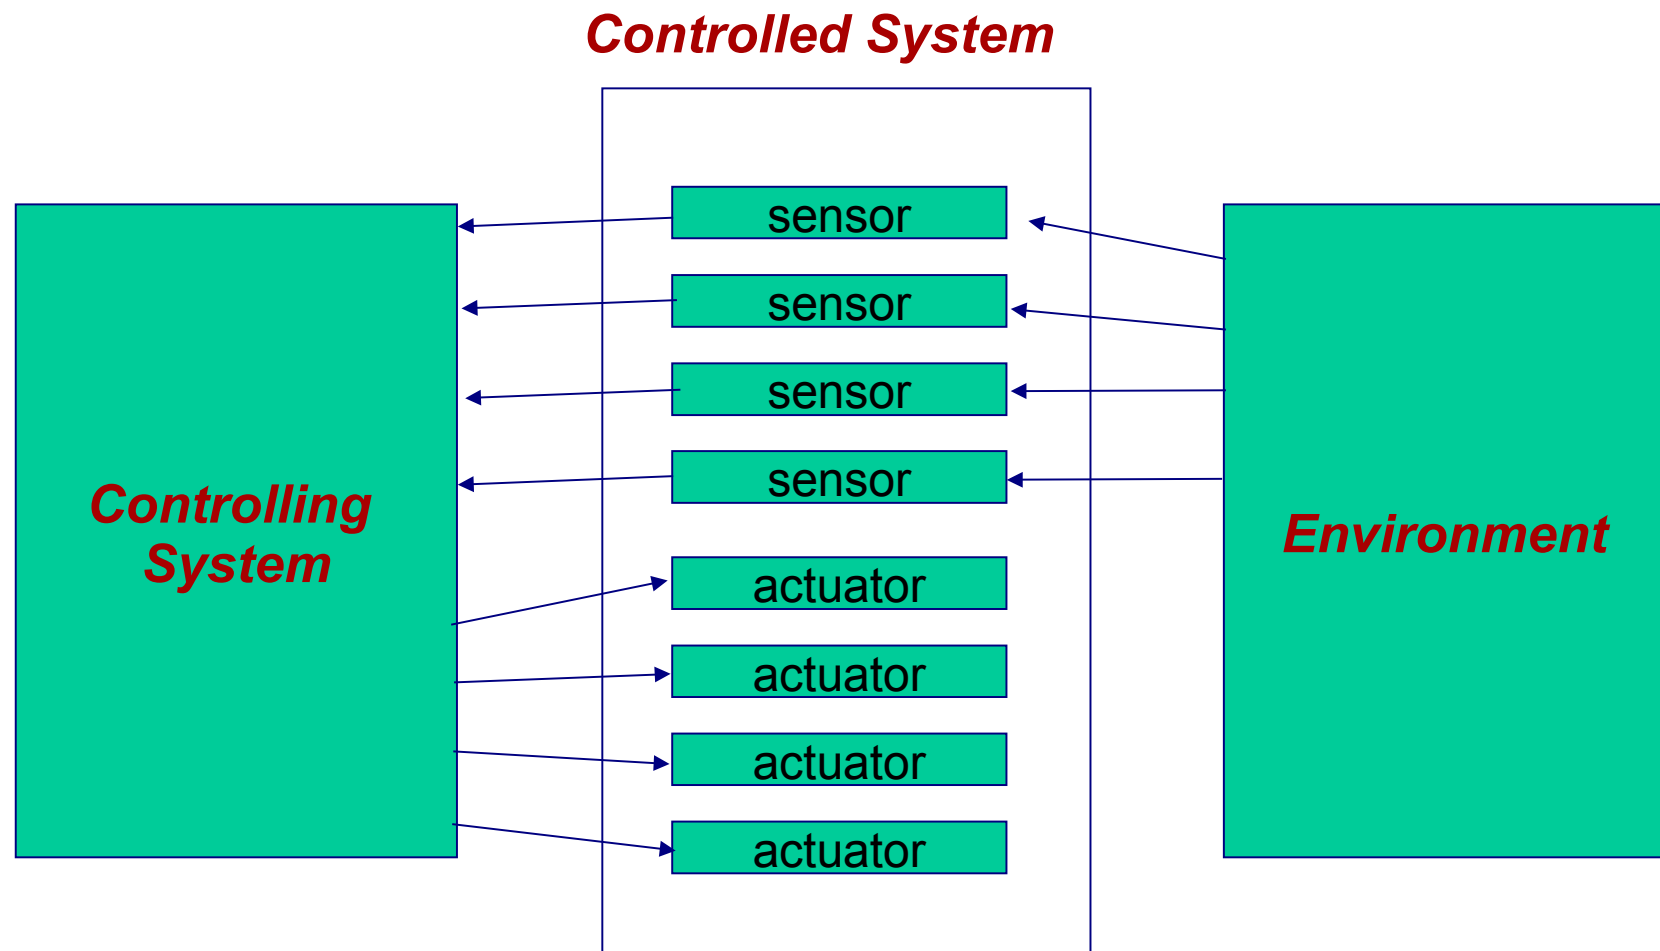
Agenda

- **Introduction to Real Time Systems (RTS)**
- **Introduction to Real Time Operating Systems (RTOS)**
 - How they differ: RTOS vs GPOS
 - RTOS Building Blocks: Tasks, Scheduler, Services
- **Overview of popular RTOS**

Real-Time System definition

- **Real-time systems have been defined as: "those systems in which the correctness of the system depends not only on the logical result of the computation, but also on the time at which the results are produced";**
 - The misconception that real time means fast is wrong
 - J. Stankovic, "Misconceptions About Real-Time Computing," IEEE Computer, 21(10), October 1988.

Typical Real-Time System



Real Time System Characteristics

- **Correctness**

- The correctness of the system depends not only on the logical result, but also on the time at which the results are produced.

- **Predictability**

- Possible to show at “design time” that all the timing constraints of the application will be met.

- **Deterministic**

- For each possible state and each set of inputs, a unique set of outputs and next state of the system can be determined.

Real Time Tasks

- **Periodic tasks**

- Time-driven (p_i , c_i)
- Eg: Task monitoring temperature of a patient in an ICU

- **Sporadic Tasks**

- Event-driven (c_i , g_i , d_i)
- Minimum separation (g_i) between two consecutive instances of the task implies that once an instance of a sporadic task occurs, the next instance cannot occur before g_i .
- Hard deadline. Eg: Reporting of fire conditions

- **Aperiodic tasks**

- Similar to Sporadic tasks
- Minimum separation between two consecutive instances can be 0
- Soft deadline Eg: Interactive commands issued by users

c_i : worst case exec time, p_i : task period, d_i : deadline, g_i : minimum separation time

Classification of Real Time Systems

- **Soft real-time system**
 - Performance is degraded but not destroyed by failure to meet response-time constraints.
 - Eg: Multimedia, Interactive video games, Online transaction systems
- **Hard real-time system**
 - Failure to meet a single deadline may lead to catastrophic failure
 - Eg: Aircraft Control Systems, Nuclear Power Stations, Chemical Plants, Life support systems
- **Synchronous real-time system**
 - Periodic (Time driven) – task activity occurs repeatedly
- **Asynchronous real-time system**
 - Aperiodic (Event driven) – task activities are asynchronous
- **Combination of both - Some tasks are periodic others, aperiodic**

Real Time Operating Systems

Goals of an Operating System

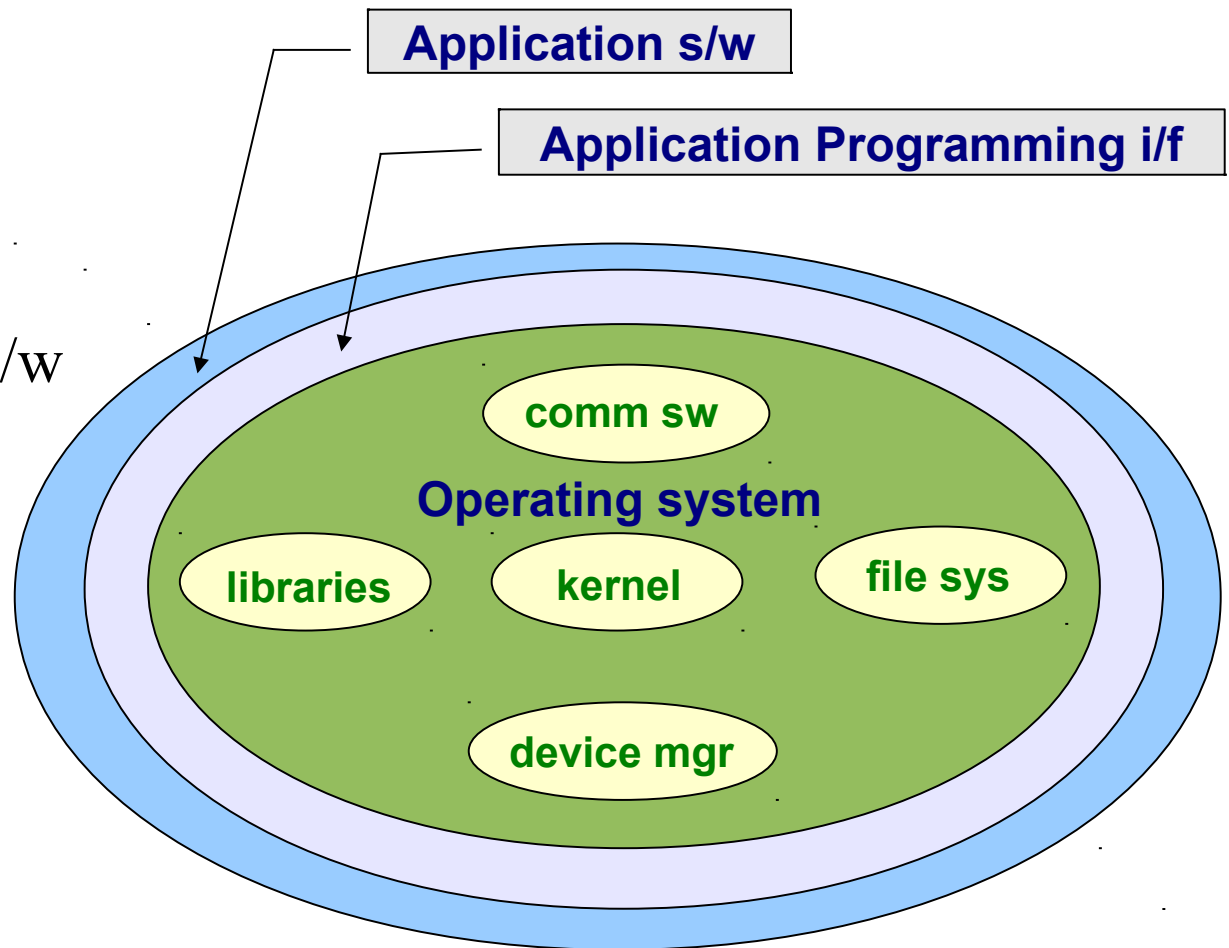
- **General Purpose Operating System**
 - Maximum Throughput and CPU Utilization
 - User experience through enhanced graphics
- **Real Time Operating System**
 - Handling timing constraints imposed by the application
 - Respond to events quickly
 - Determinism

RTOS Vs GPOS

#	Evaluation Metrics	General Purpose OS	Real Time OS
1	Determinism	<ul style="list-style-type: none">• Non Deterministic	<ul style="list-style-type: none">• All RTOS functions should execute in a fixed/deterministic amount of time
2	Load Independent Timing	<ul style="list-style-type: none">• Not Applicable• Response becomes sluggish as number of tasks increase	<ul style="list-style-type: none">• Remains Constant• Irrespective of system load, performance of the RT system should remain predictable
3	Task Level Scheduling	<ul style="list-style-type: none">• Generally Round Robin Scheduling• Sometimes Priority based scheduling• Efforts are made to ensure that all tasks get a chance to execute	<ul style="list-style-type: none">• Generally Priority based Preemptive Scheduling .• Time slicing only among tasks that hold the same priority• Ensure that the Highest Priority task is always executed, even if it is the most frequent
4	Interrupt Management	<ul style="list-style-type: none">• Nesting may be disabled.• Interrupt latency, response and recovery are not performance metrics	<ul style="list-style-type: none">• Nesting is always enabled.• Interrupt latency, response and recovery are very important performance metrics

RTOS Architecture

- An RTOS consists of
 - Kernel
 - Device Manager
 - Networking protocol s/w
 - Libraries
 - File system (optional)



RTOS Services

- **Multitasking, Task & Thread Management - creation and scheduling of tasks, priority based preemptive scheduling**
- **Vectored Interrupt Service Routines - avoid polling**
- **Inter Task Synchronization and Inter Task Communications through mutual exclusions, signals, messages, shared memory, etc**
- **Timer Services such as periodic and aperiodic interrupts**

RTOS Services

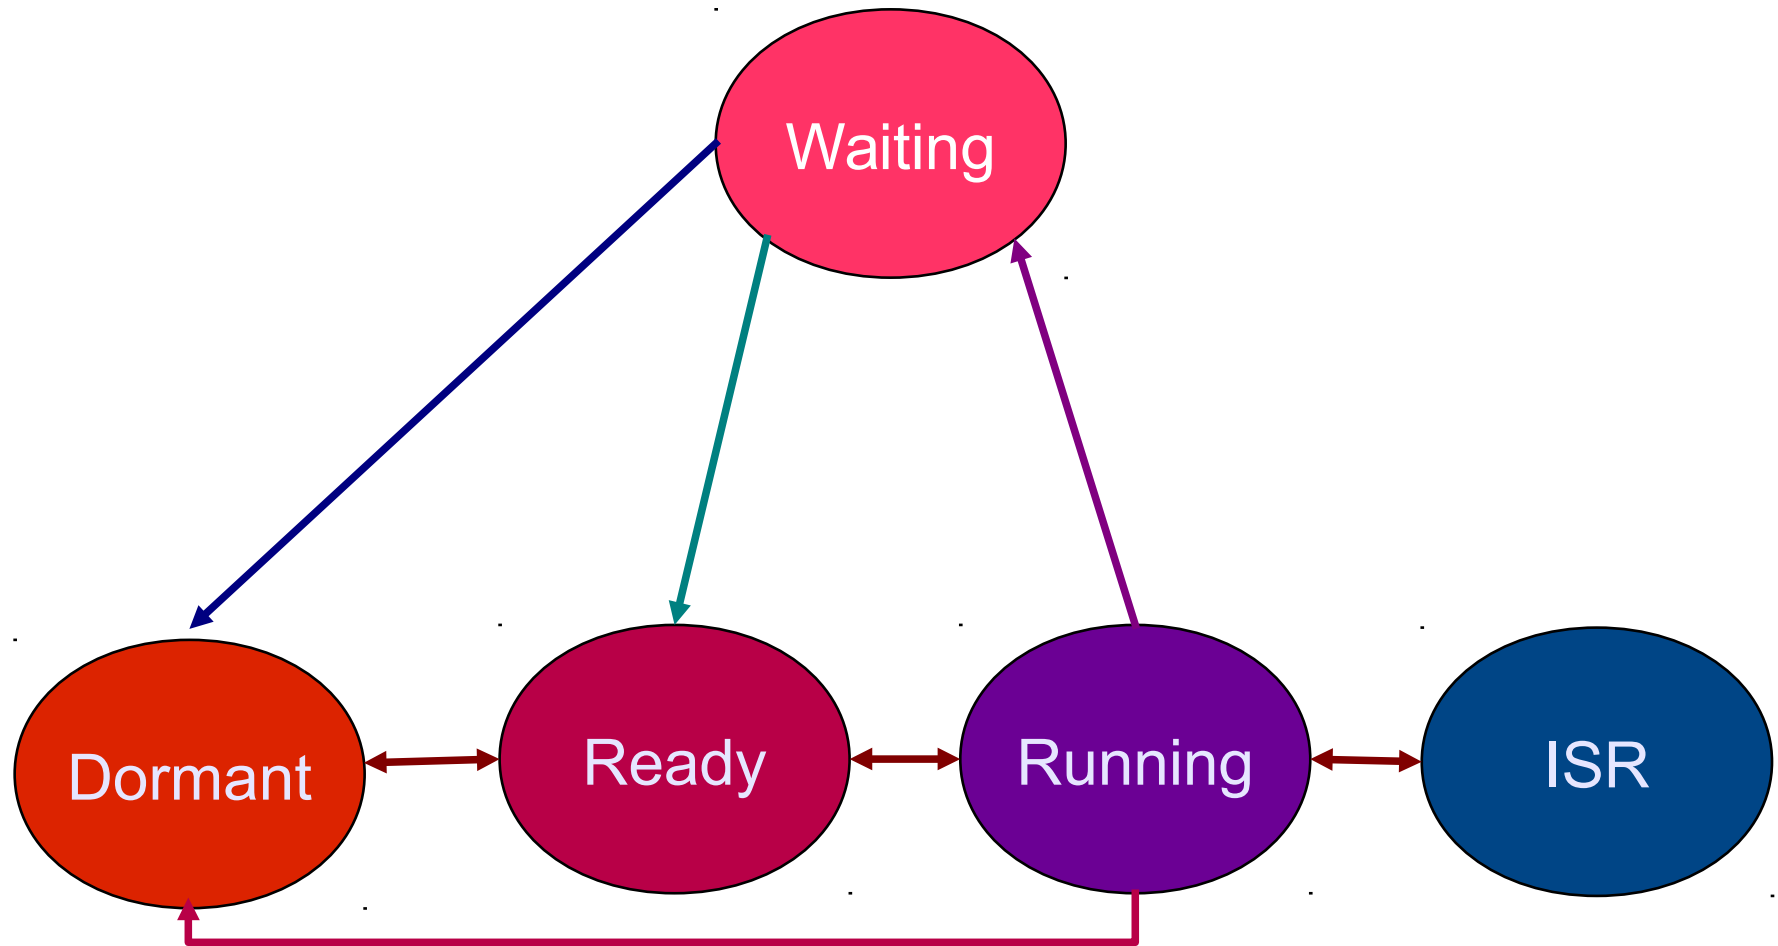
- **Device drivers to service the needed special devices**
- **Communication Protocol software (networking through Ethernet, wireless, etc)**
- **Application Programming Interface (APIs) to access kernel Services**

Real Time Tasks & Real Time Scheduling

Real Time Tasks

- **An RT application splits the work into Tasks.**
- **A real time task, also called a thread, is a simple program that thinks it has the CPU all to itself.**
- **It is generally implemented as an infinite loop (periodic tasks).**
- **Each Task is assigned its own set of CPU Registers, Stack Area and a Priority.**
- **A Task may be visualized through 3 logical components**
 - **The Task Function** (Logic) – What the task is out to achieve
 - **The Task Control Block** – Holds configuration information of the task like Priority, State, Period, Stack Pointer
 - **The Task Stack** – Maintains context information of the task during switches

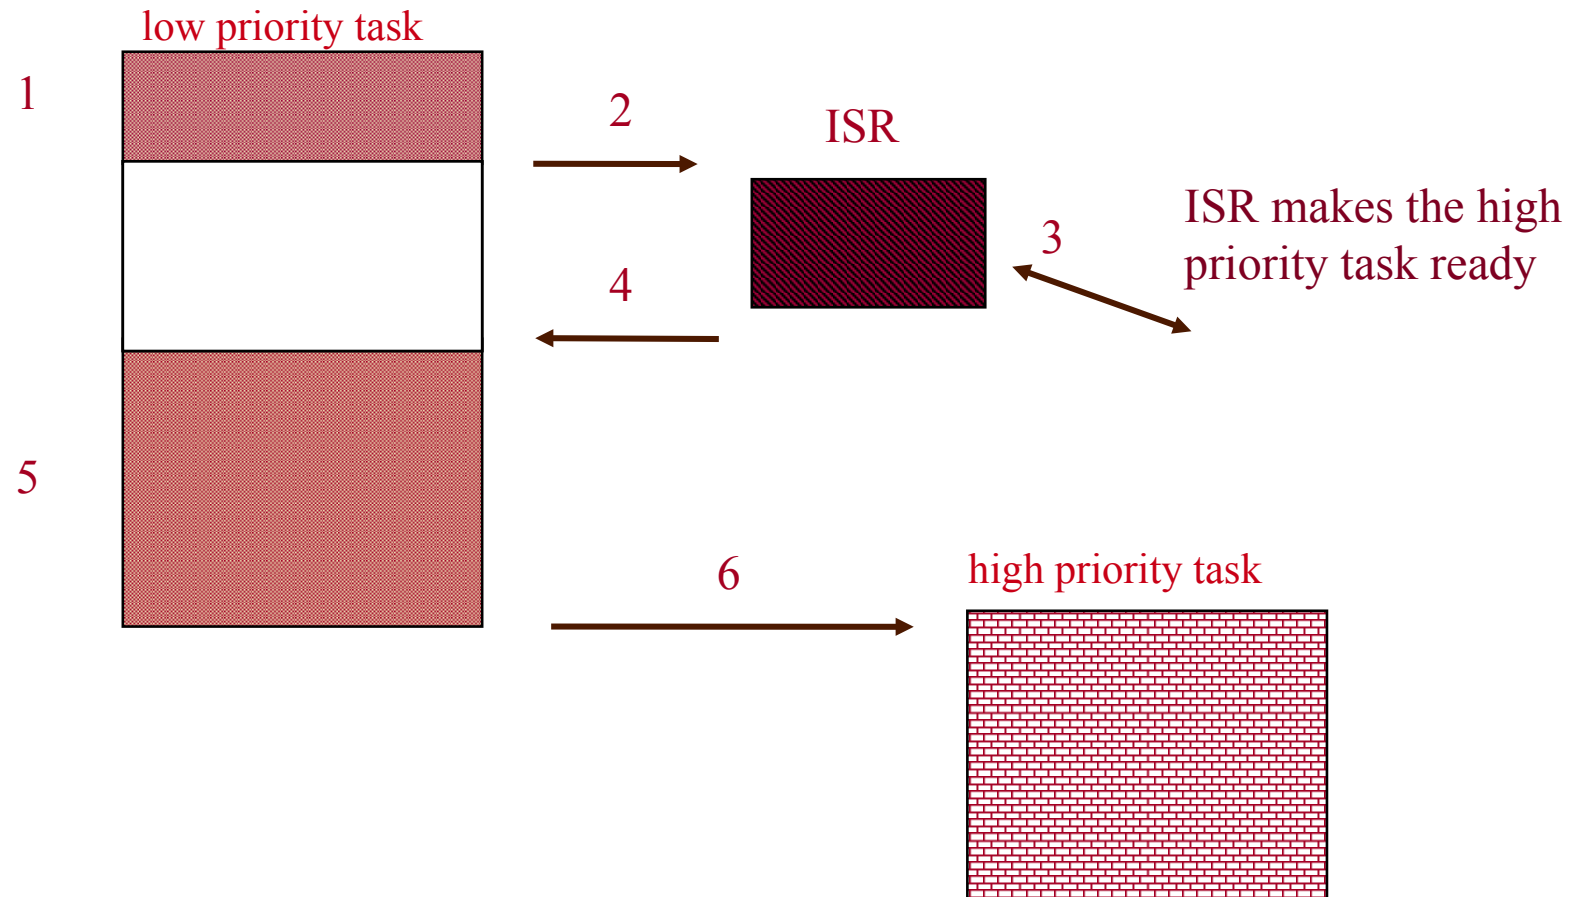
Task States



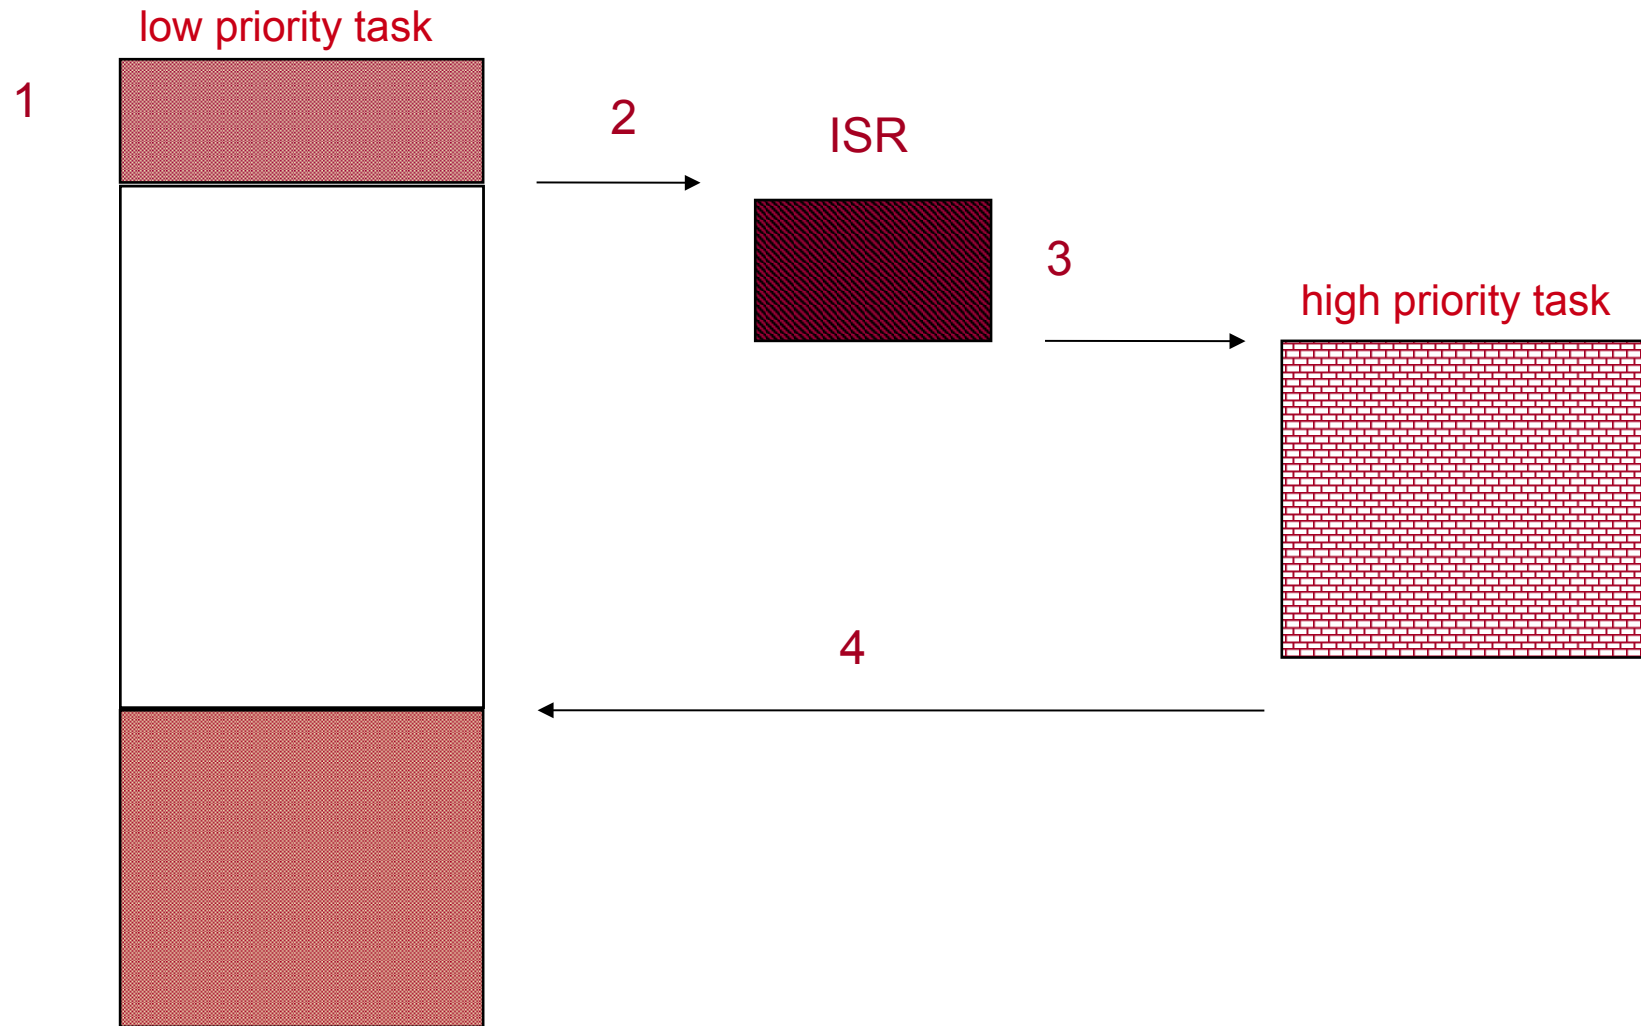
Real-time Scheduling

- **Goals**
 - Meeting the timing constraints of the system
 - Preventing simultaneous access to shared resources and devices
 - Attaining a high degree of CPU utilization while satisfying the timing constraints of the system
 - Reducing the cost of context switches caused by preemption

Kernel Types - Non Preemptive Kernel



Kernel Types - Preemptive Kernel

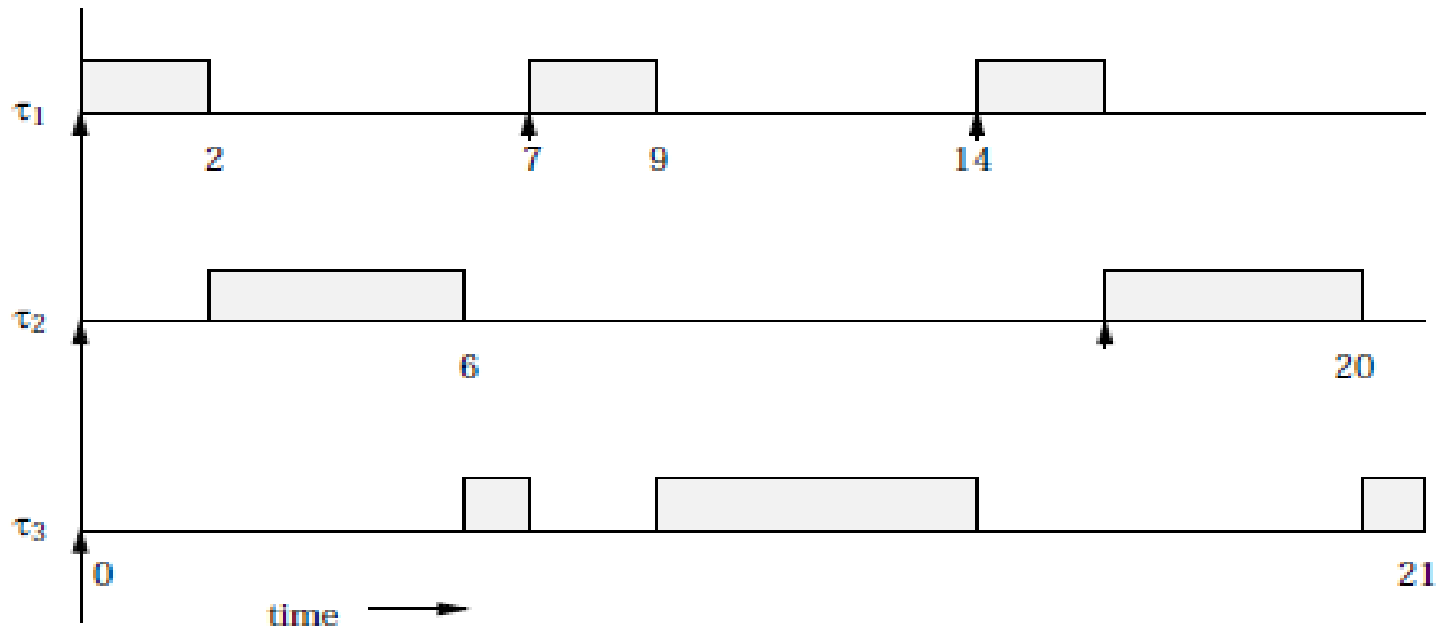


Task Timing Properties

- **Each task(T_i) occurring in a real-time system has some timing properties, which should be considered when scheduling tasks.**
 - Release time (or ready time)
 - Deadline(D_i)
 - Worst case execution time(C_i)
 - Period(T_i or P_i) (For periodic tasks)

Schedulability analysis

- **Timing diagrams provide a good way to visualize and even to calculate the timing properties of simple programs.**



Schedulability analysis

- A better method of analysis is to derive conditions to be satisfied by the timing properties of a program, to meet its deadlines
- Necessary Conditions:
 - Worst case execution time must be less than period ($C_i < T_i$)
 - CPU Utilization factor U_i for a periodic task T_i is (C_i/P_i)
 - Overall system utilization (U) = $\sum U_i \leq 1$

Utilization (%)	Zone Type	Typical Application
0-25	significant excess processing power – CPU may be more powerful than necessary	various
26-50	very safe	various
51-68	safe	various
69	theoretical limit	embedded systems
70-82	questionable	embedded systems
83-99	dangerous	embedded systems
100+	overload	stressed systems

RM, EDF, LLF Scheduling

Rate Monotonic

- Static-priority preemptive scheme
- Assumes that all tasks are periodic
- The priorities are inversely proportional to the period of the task.

Earliest Deadline First

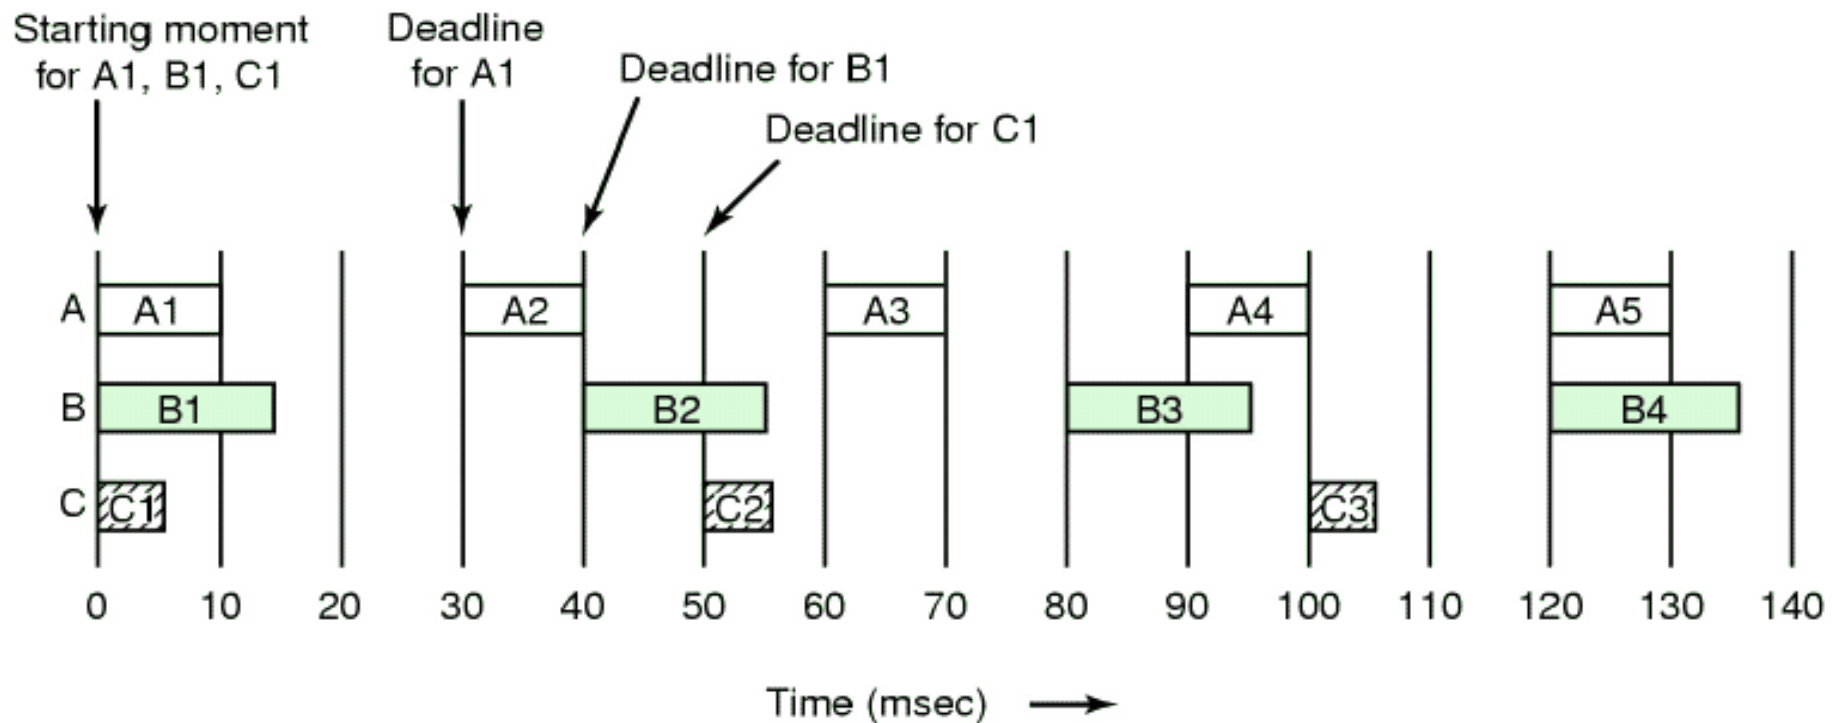
- Dynamic-priority preemptive scheme
- Higher priority is assigned to the task that has earlier deadline

Least Laxity First

- Dynamic priority preemptive scheme
- The laxity of a task is deadline minus remaining computation time.
- Highest priority is given to the task with least laxity

Scheduling Example

- Consider 3 tasks with the following timing properties



Scheduling Example

- Are the tasks schedulable ??

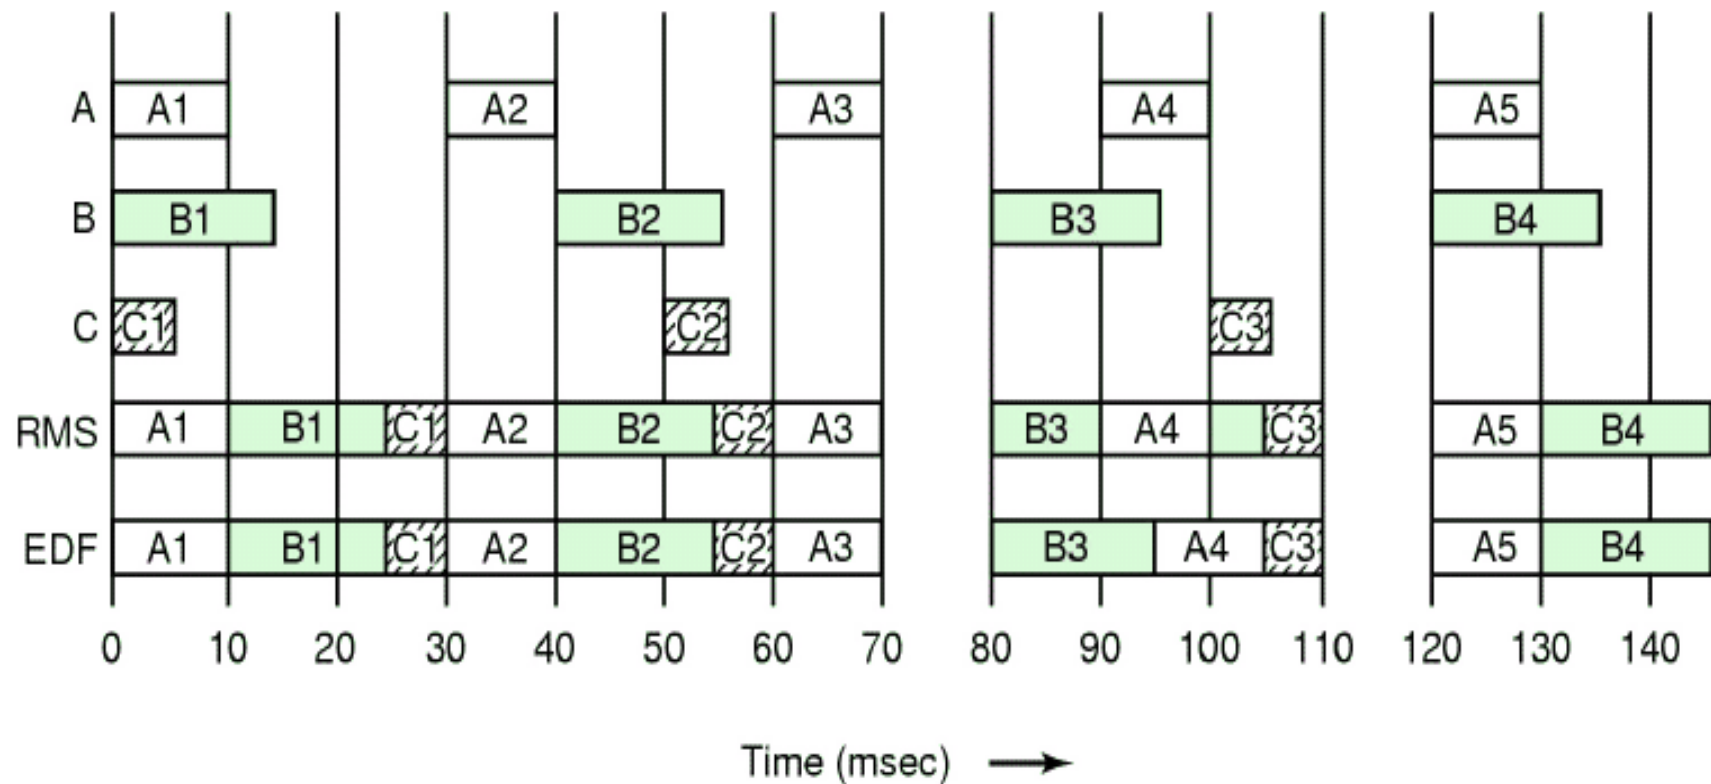
$$\sum_{i=1}^m \frac{C_i}{P_i} \leq 1$$

$$\frac{10}{30} + \frac{15}{40} + \frac{5}{50} = 0.808$$

- YES

Scheduling with RM and EDF

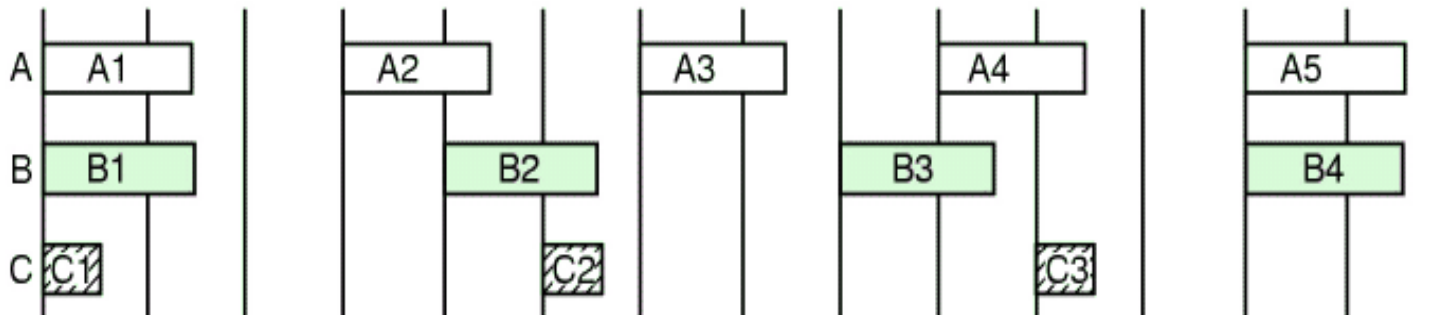
- Analyze using timing diagram



Modify Example

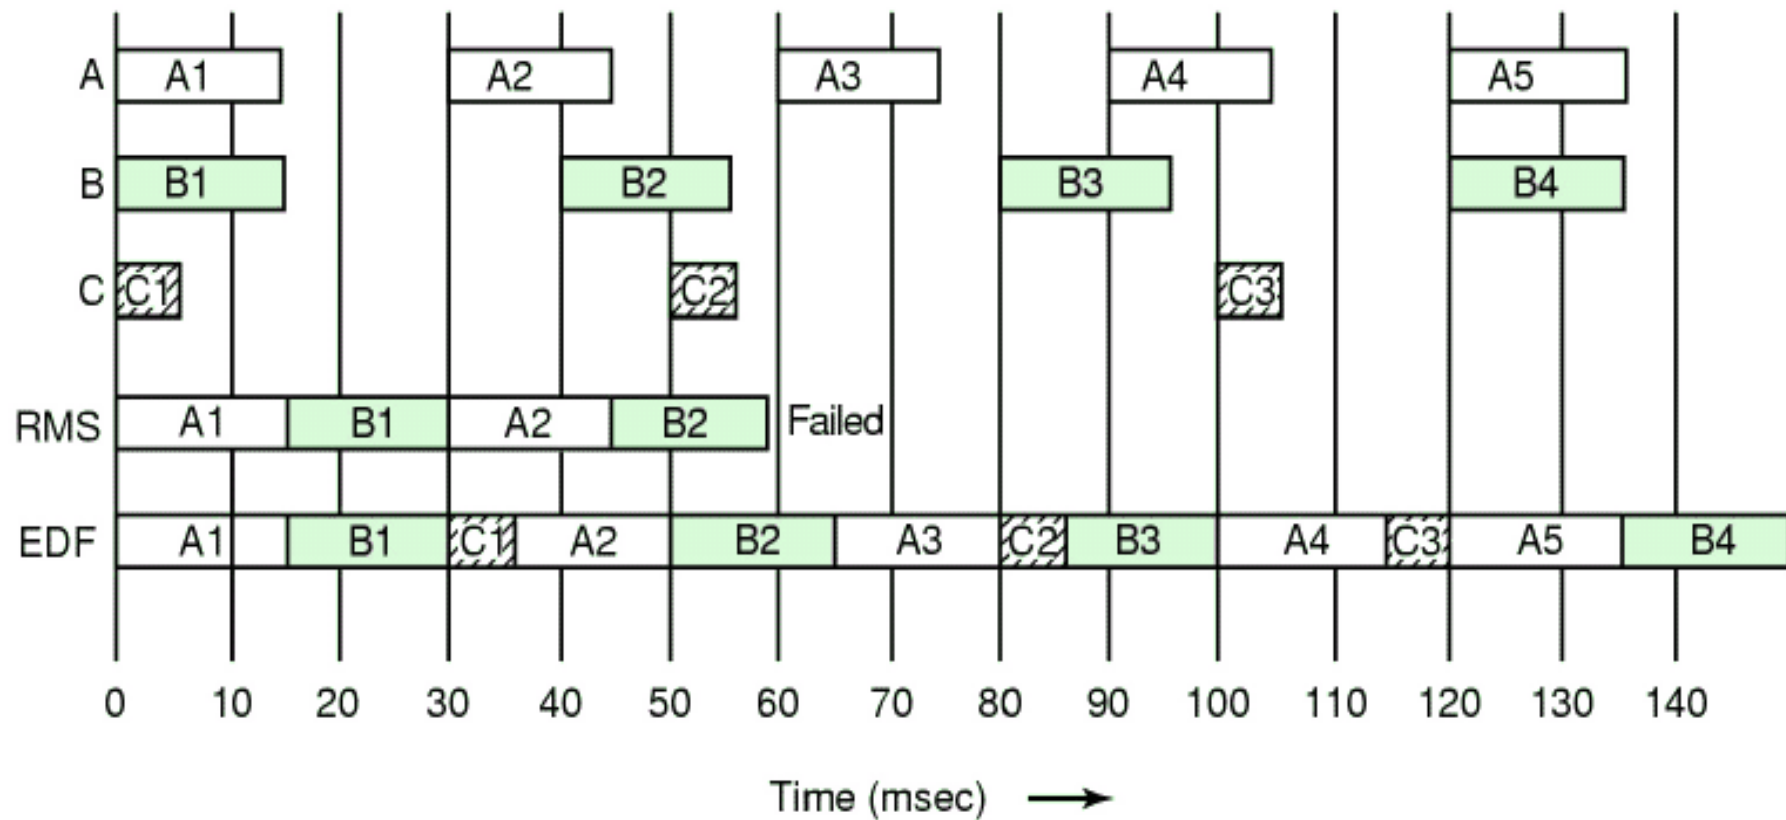
- Increase A's CPU requirement to 15 milli sec
- The system is still schedulable

$$\frac{15}{30} + \frac{15}{40} + \frac{5}{50} = 0.975$$



Scheduling using RM and EDF

- Analyze using timing diagram



RM Scheduling Condition

- **RM Scheduling is guaranteed to work if the CPU utilization is not too high**
- **Condition of Schedulability**

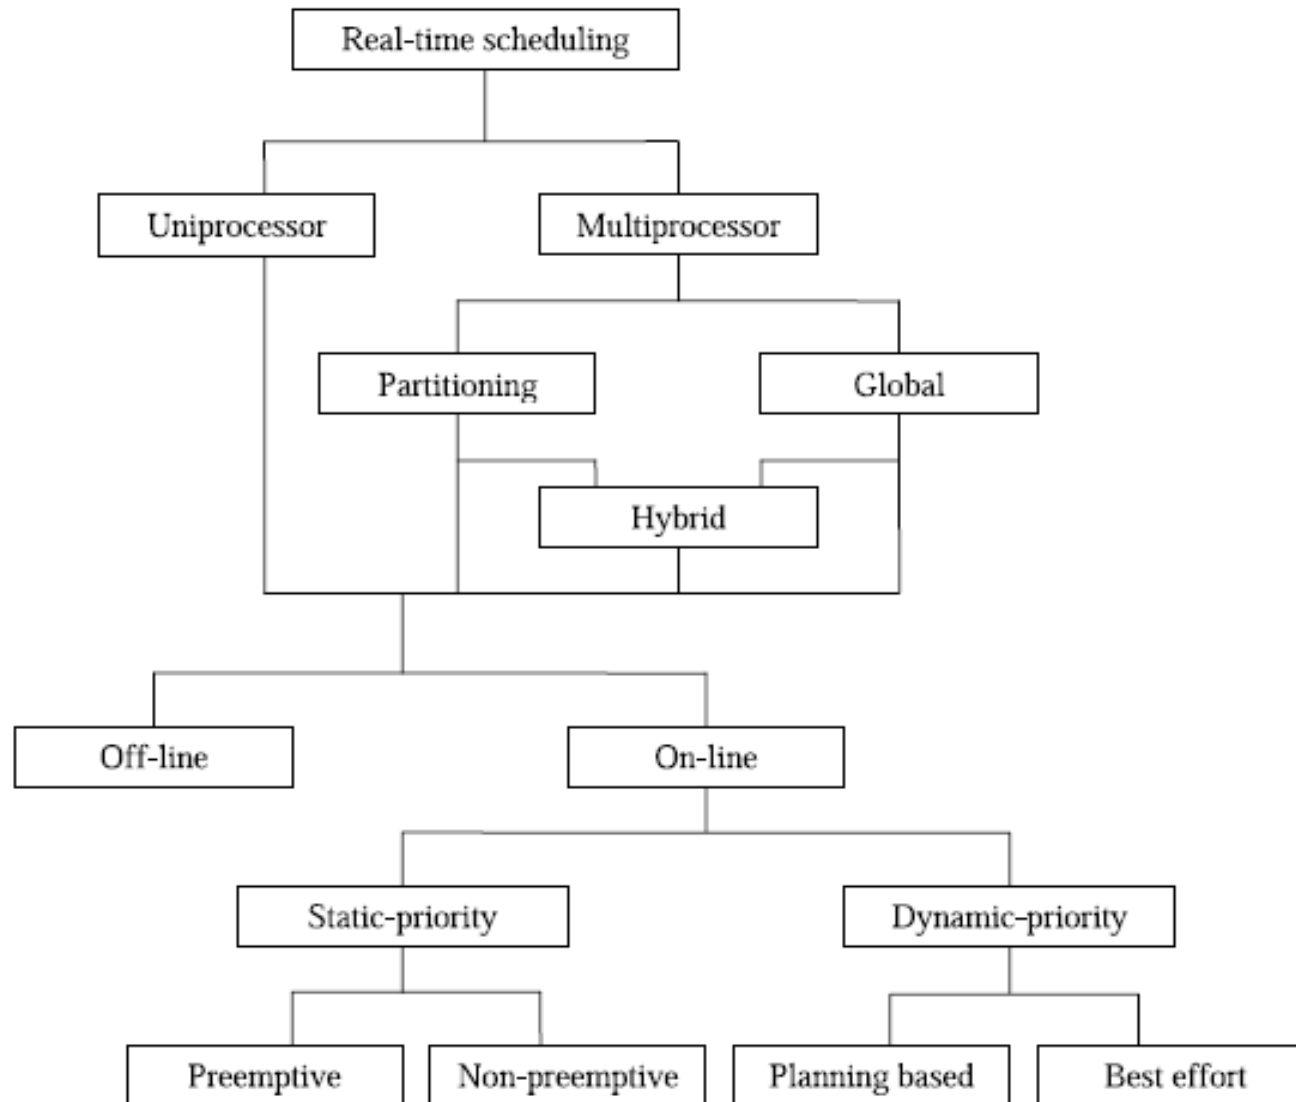
$$\sum_{i=1}^m \frac{C_i}{P_i} \leq m(2^{1/m} - 1)$$

- Where m = number of tasks
- For three tasks, CPU utilization must be less than 0.780

Performance Comparison

- **RM Scheduling** which is simple and easy to implement can be used for systems with low CPU utilization
- **EDF Scheduling** always works for any schedulable set of tasks. Upto 100% CPU utilization
- **LLF Scheduling** is similar in properties to EDF
 - Takes into account that laxity time is more meaningful than deadline for tasks with mixed computing sizes

Classification of Real-time Scheduling



Multiprocessor Scheduling Algorithms

Global Scheduling Algorithms

- **Store the tasks that have arrived in a queue, which is shared among all processors**
- **Each processor maintains**
 - Status table of tasks it has committed to run.
 - Table of the surplus computational capacity at every other processor. (Each processor regularly sends to others the fraction of the next window that is free)
- **A Overloaded processor selects a processor that is most likely to be able to successfully execute that task by its deadline and ships the tasks out**
- **Eg: Focused addressing and bidding algorithm**

Multiprocessor Scheduling Algorithms

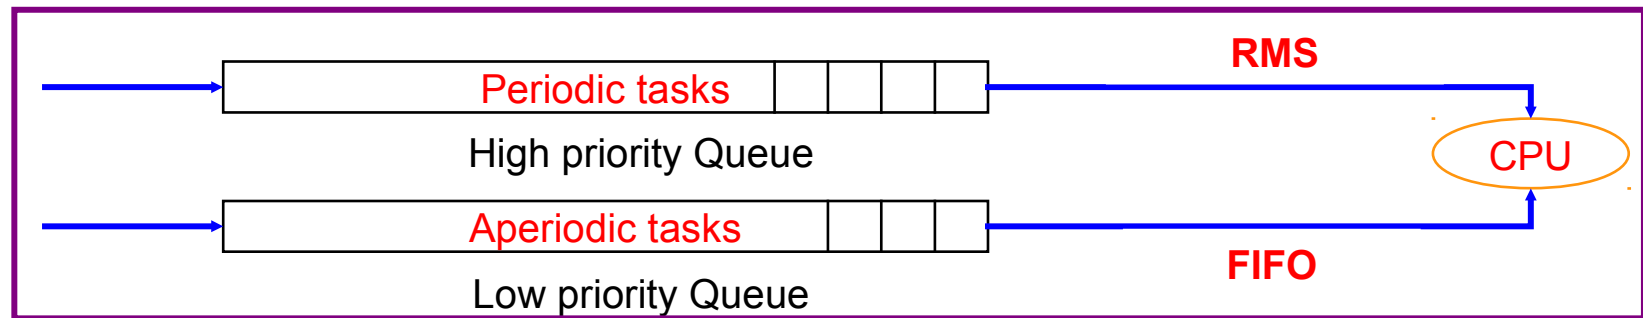
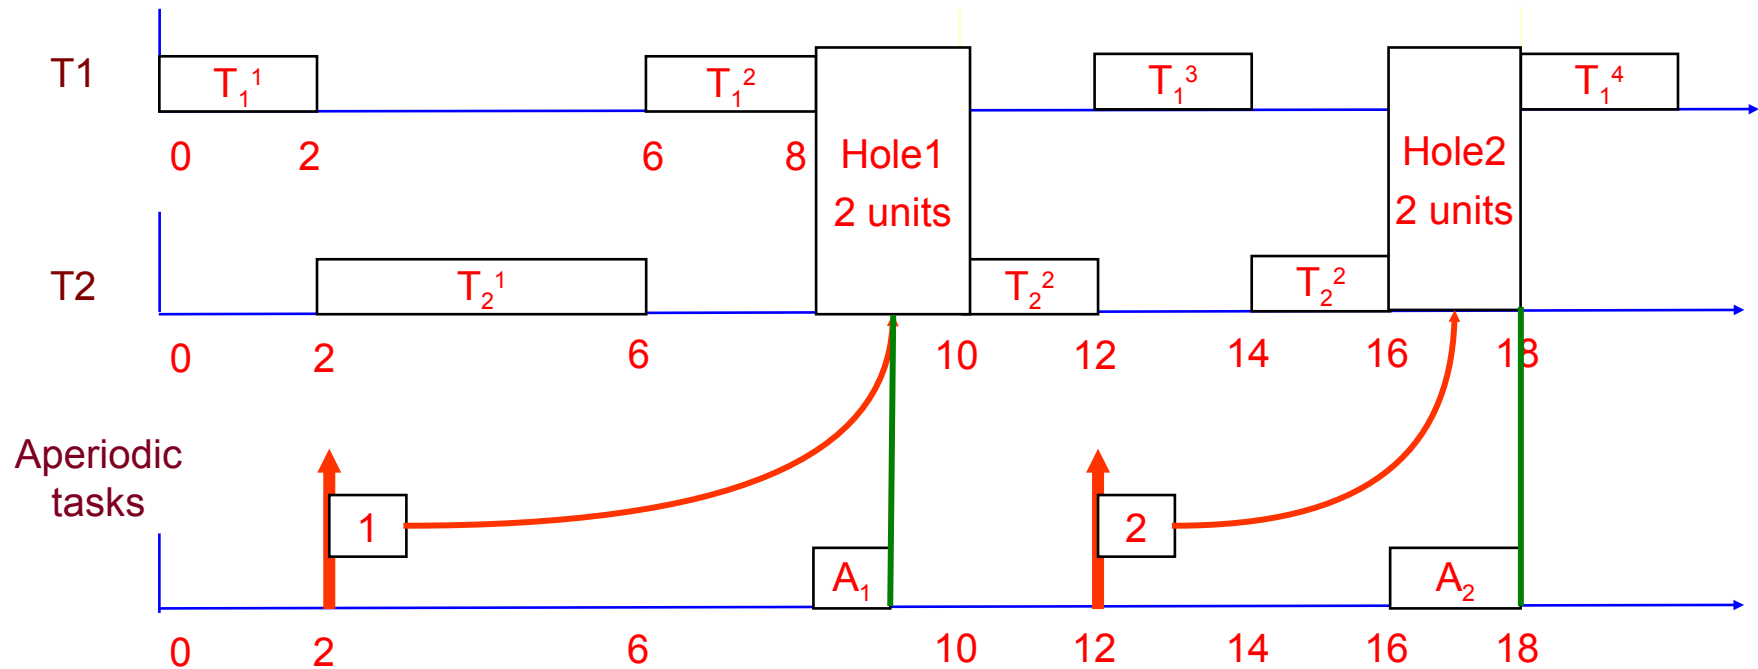
Partitioning Scheduling Algorithms

- **Tasks are partitioned such that all tasks in a partition are assigned to the same processor**
 - Tasks of same class, are guaranteed to satisfy the RM schedulability on one processor
 - Tasks are not allowed to migrate
- **Partitioning scheduling has a low scheduling overhead compared to global scheduling, because tasks do not need to migrate across processors**
- **Eg: Next fit algorithm for RM scheduling**

Scheduling of Aperiodic Tasks

- **Background:** scheduled when processor is idle
- **Interrupt-driven:** scheduled on arrival
- **Periodic server:** defined by (ps, es) for processing aperiodic tasks. Budget replenishes at ps intervals.
 - Bandwidth non-preserving server
 - If scheduled and queue empty then budget set to 0.
 - Polling server
 - Bandwidth preserving server
 - Improves on the polling server by preserving budget (bandwidth) when aperiodic queue is empty.
 - Deferrable servers
 - Priority Exchange and Deferred Servers

Background Scheduling: Example

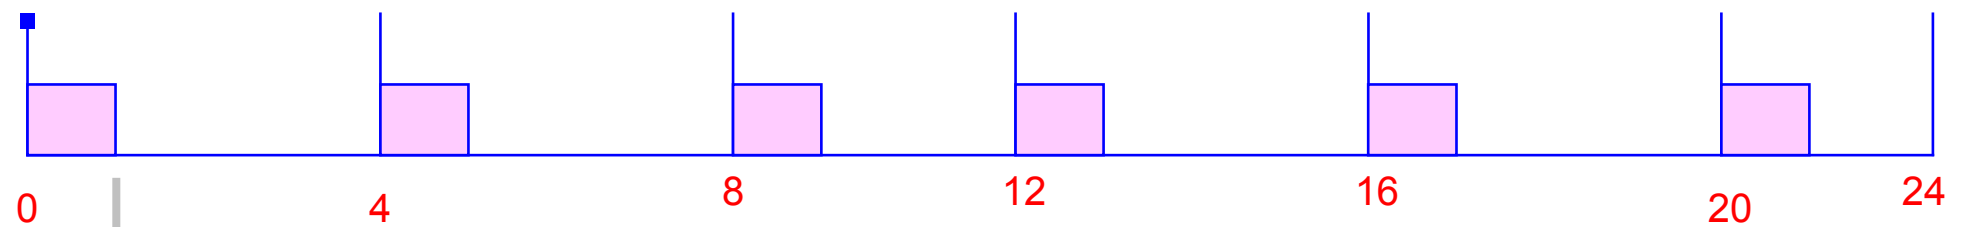


Polling server: Example

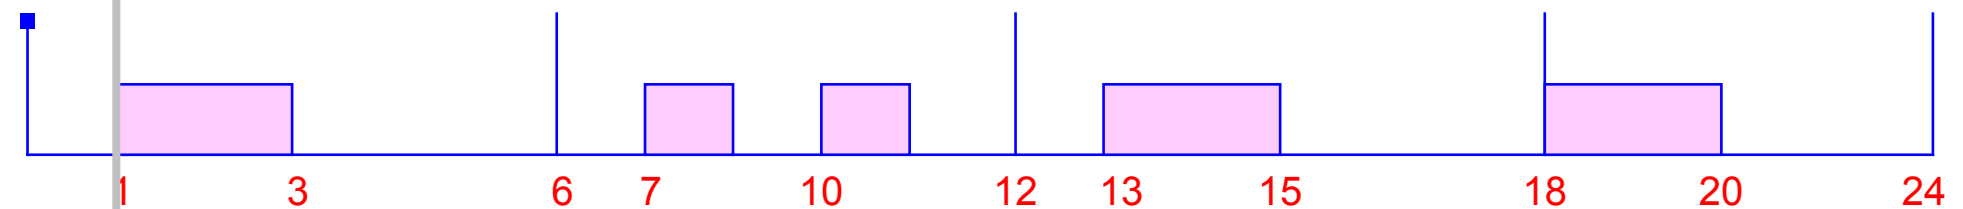
Task set: $T_i = (c_i, p_i)$

$T1 = (1,4)$, $T2 = (2,6)$ and $Ts = (2,5)$

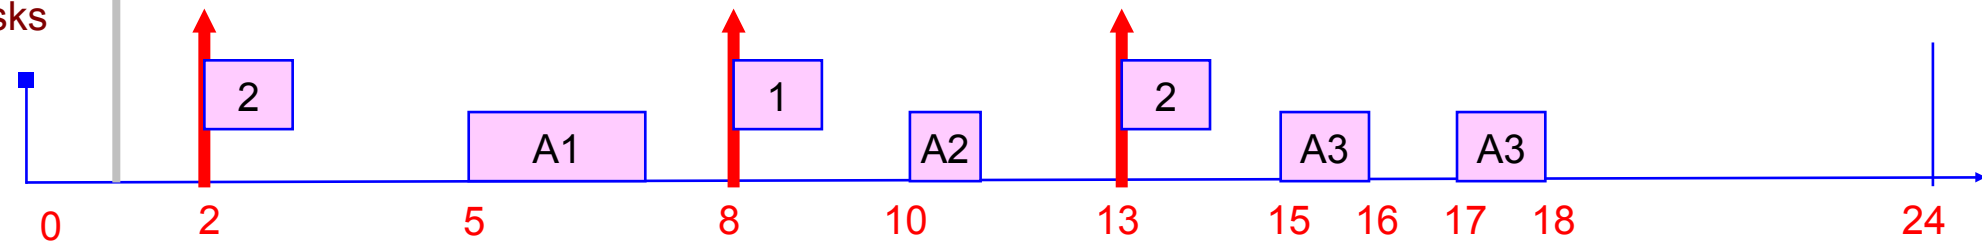
T1



T2



Aperiodic
tasks

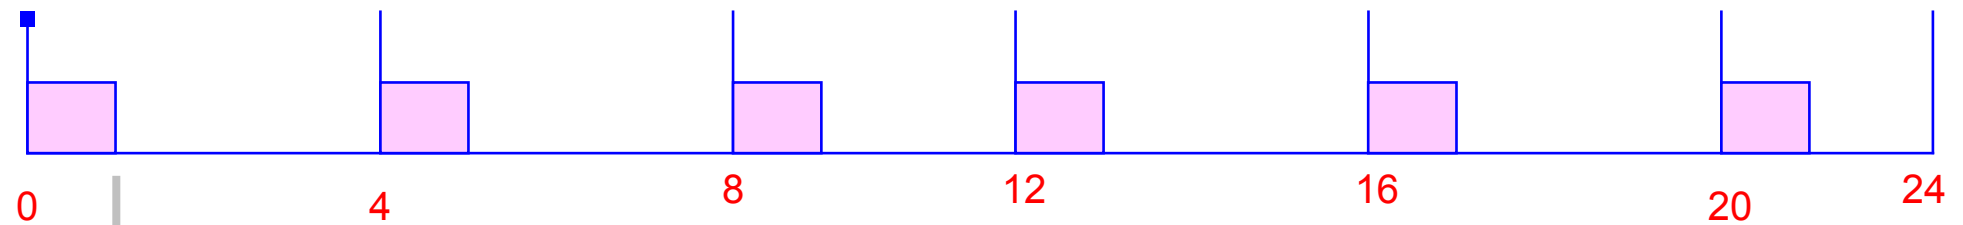


Deferred server: Example

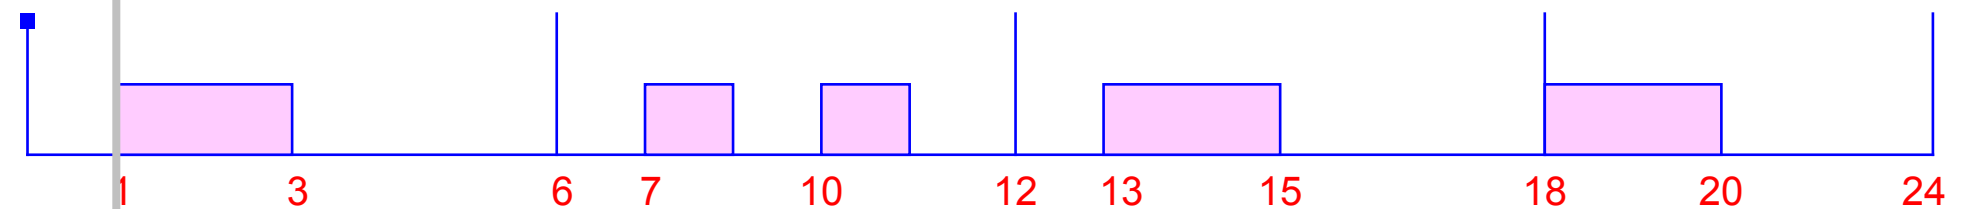
Task set: $T_i = (c_i, p_i)$

$T1 = (1,4)$, $T2 = (2,6)$ and $Ts = (2,5)$

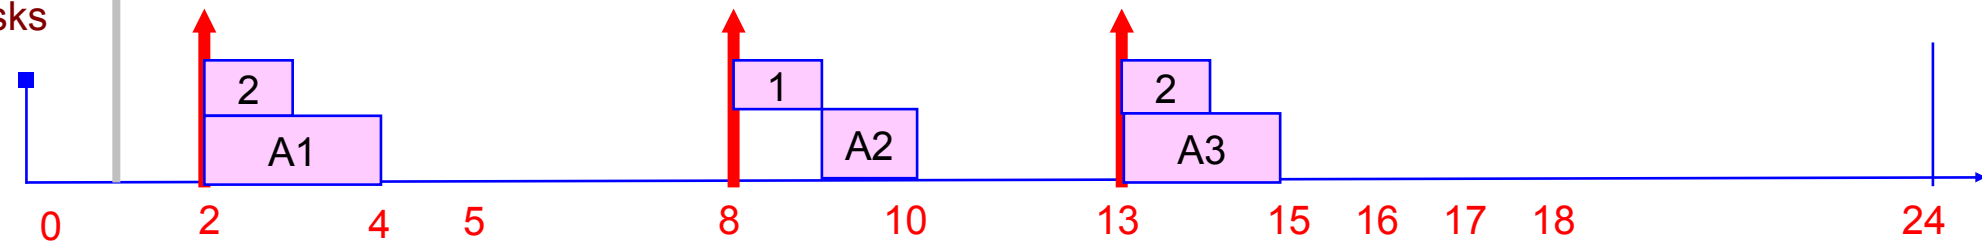
T1



T2



Aperiodic tasks



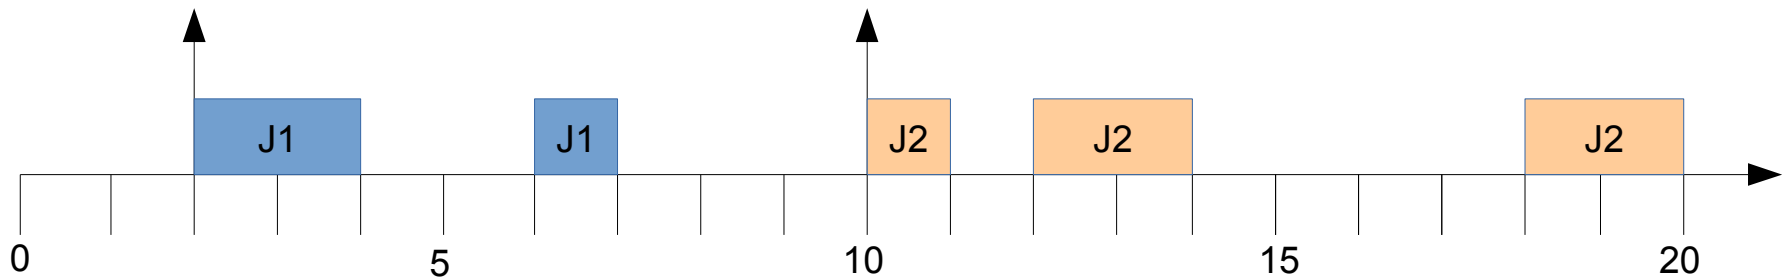
Scheduling of Sporadic Tasks

- **Methods**

- Consider sporadic tasks as periodic tasks with a period equal to their minimum inter-arrival time
- Define a fictitious periodic task of highest priority and when the task is scheduled, run any sporadic task awaiting service
- The Deferred server
 - When sporadic tasks are scheduled and none is awaiting service, it schedules periodic tasks in order of priority
 - Wastes less bandwidth

Scheduling of Sporadic Tasks

- **Consider a deferred server $js(2, 6)$ for a collection of sporadic tasks**
 - Sporadic task $j1$ with $c1 = 3$ arrives at time 2
 - Sporadic task $j2$ with $c2 = 5$ arrives at time 10

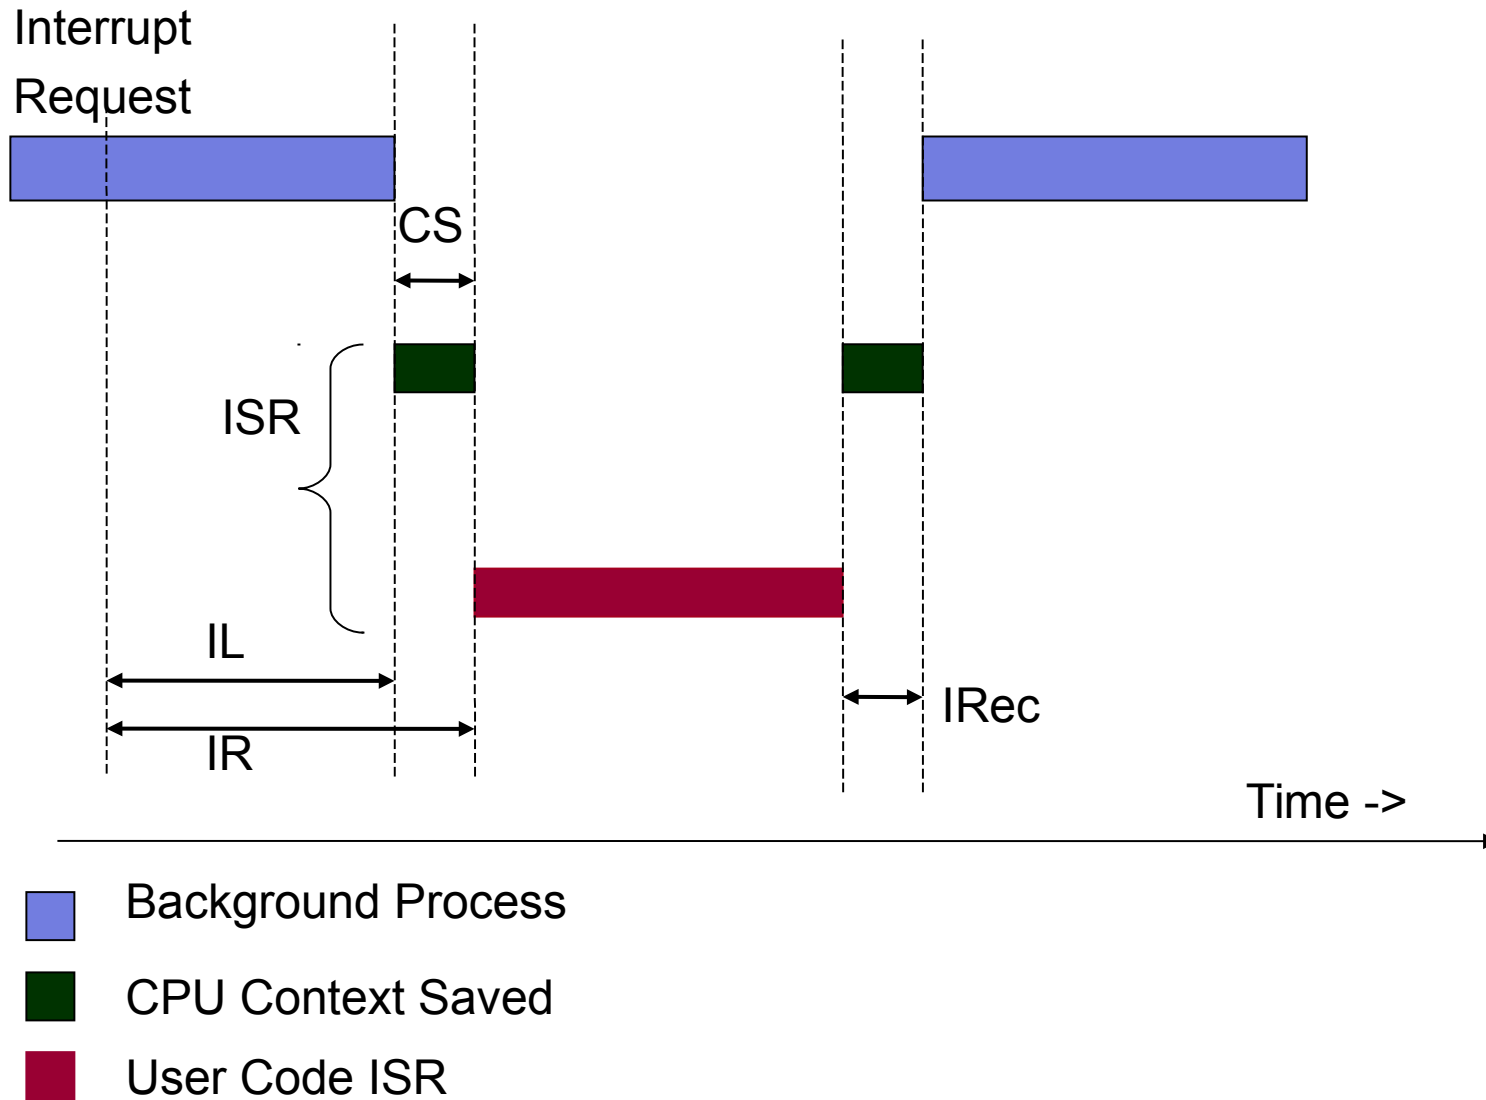


Interrupt Service Routines

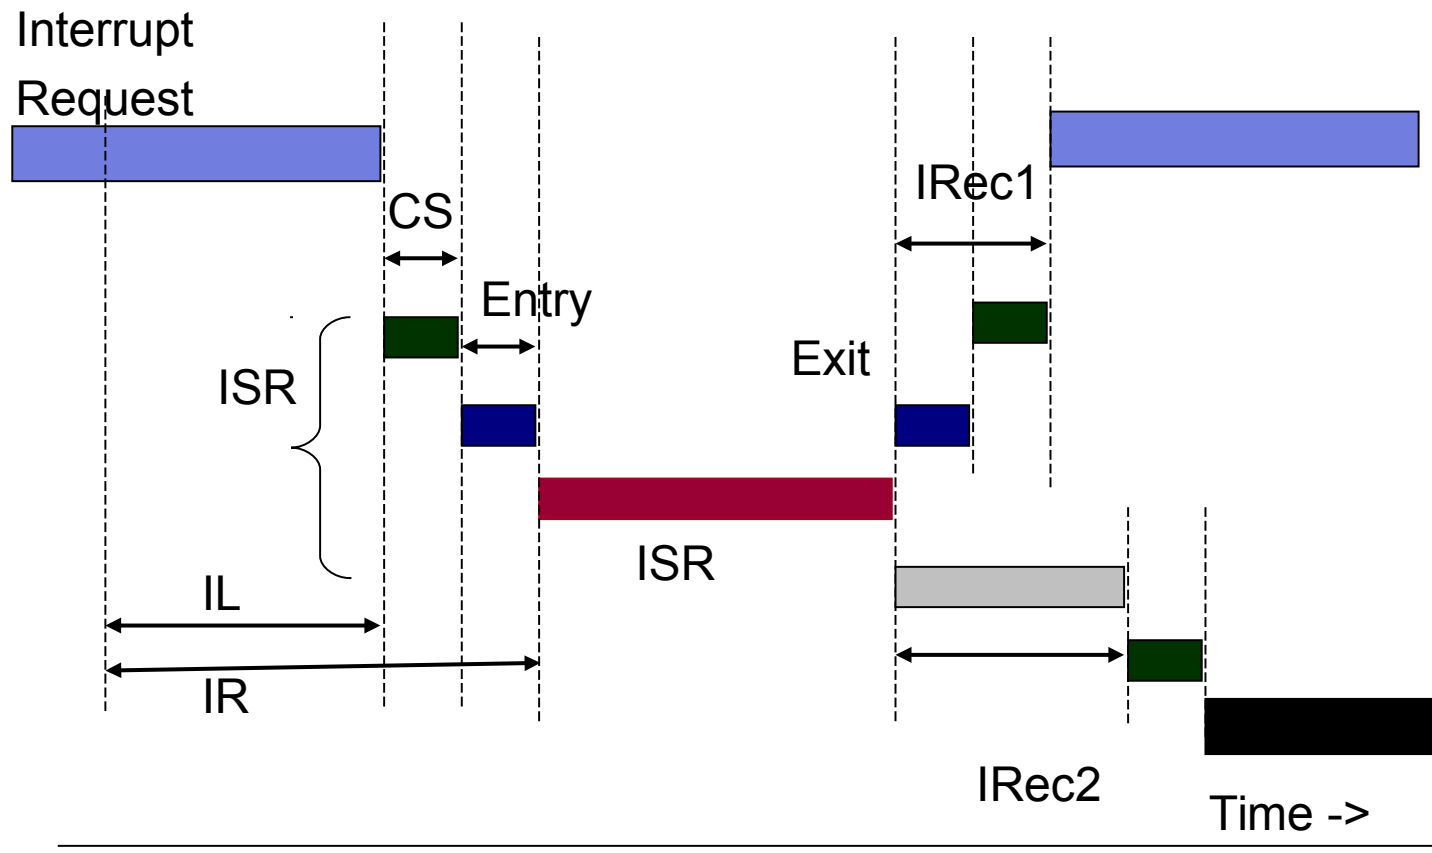
Interrupts

- **When an interrupt occurs**
 - CPU saves its context on the stack, Jumps to the Interrupt Servicing Routine (ISR), Executes ISR and Returns
 - **Interrupt Latency**
 - max time interrupts are disabled + time to begin servicing the interrupt
 - **Interrupt Response Time**
 - Interrupt Latency + time to start execution of 1st instruction in ISR
 - **Interrupt Recovery Time**
 - time for CPU to return to interrupted code / highest priority task

Interrupts in Non Preemptive Kernels



Interrupts in a preemptive kernel

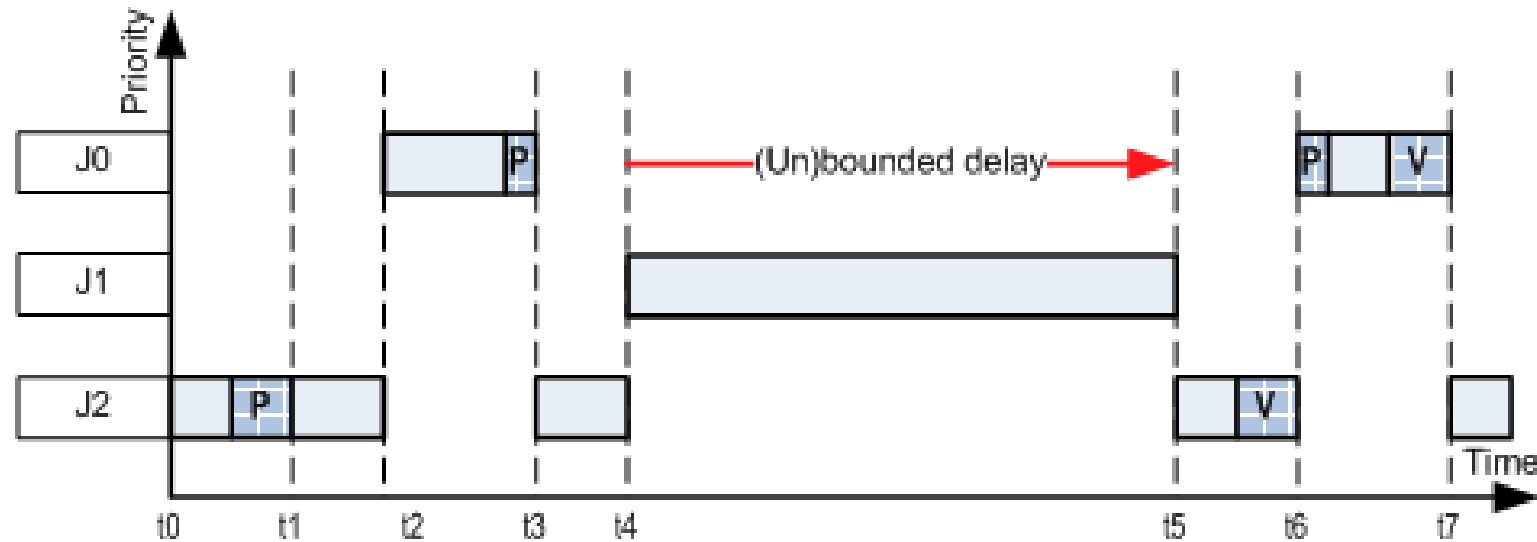


- Background Process
- CPU Context Saved
- User Code ISR

Preemptive Kernel

Inter Task Synchronization

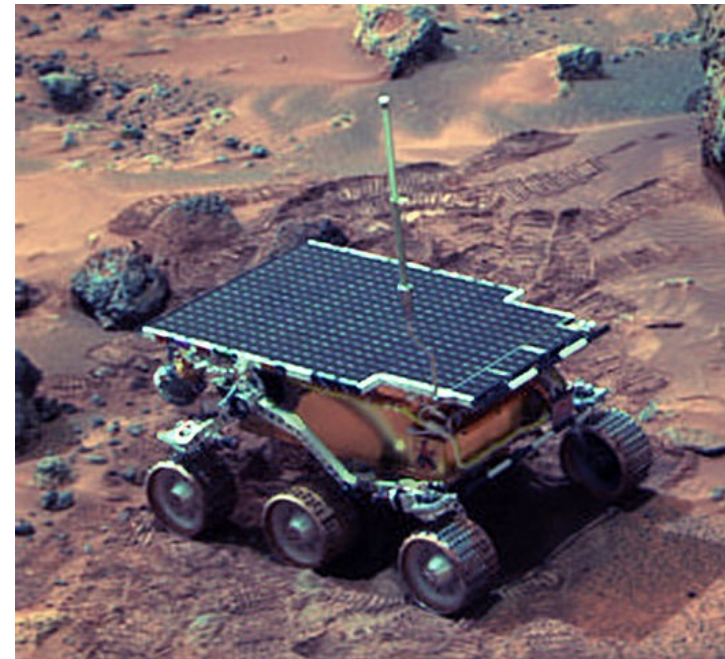
Priority Inversion Problem



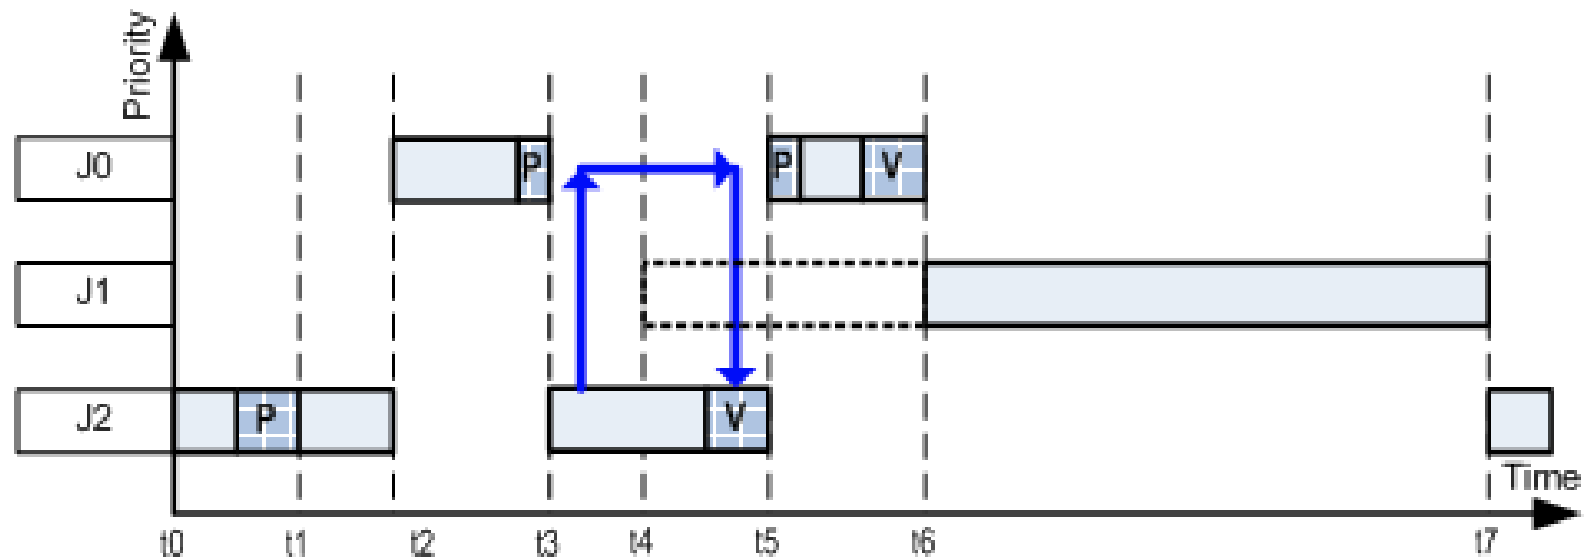
- **Effectively, the high priority task, J0, is blocked by medium priority task J1. This condition is called priority inversion.**
- **Primary techniques to solve priority-inversion**
 - Priority Inheritance Protocol
 - Priority Ceiling Protocol

Mars Pathfinder

- **The flow of information and images was interrupted by a series of system resets**
 - The Pathfinder's applications were scheduled by the VxWorks RTOS using pre-emptive priority scheduling
 - The meteorological data gathering task ran as an infrequent, low priority thread, and used the information bus synchronized with mutual exclusion locks
 - A very high priority bus management task, also accessed the bus with the same mutexes
 - A long-running communications task, having higher priority than the meteorological task, but lower than the bus management task, prevented it from running.
 - Watchdog timer noticed that the bus management task had not been executed for some time, concluded that something had gone wrong, and ordered a total system reset

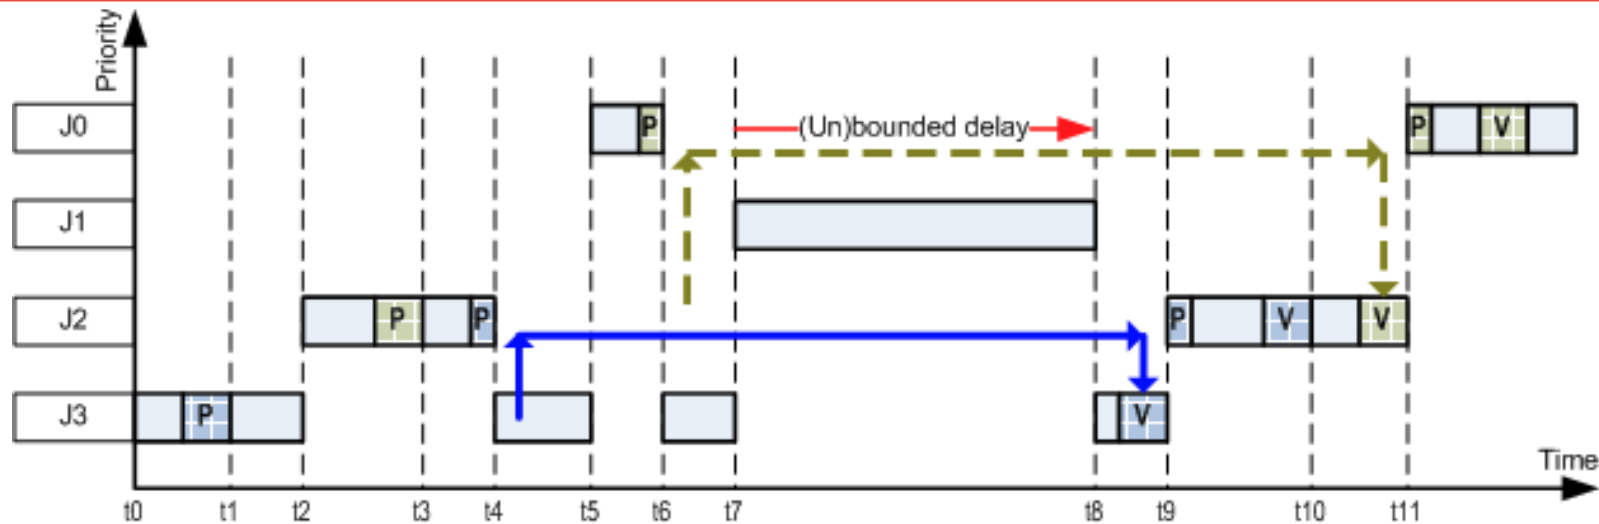


Simple Priority Inheritance Protocol

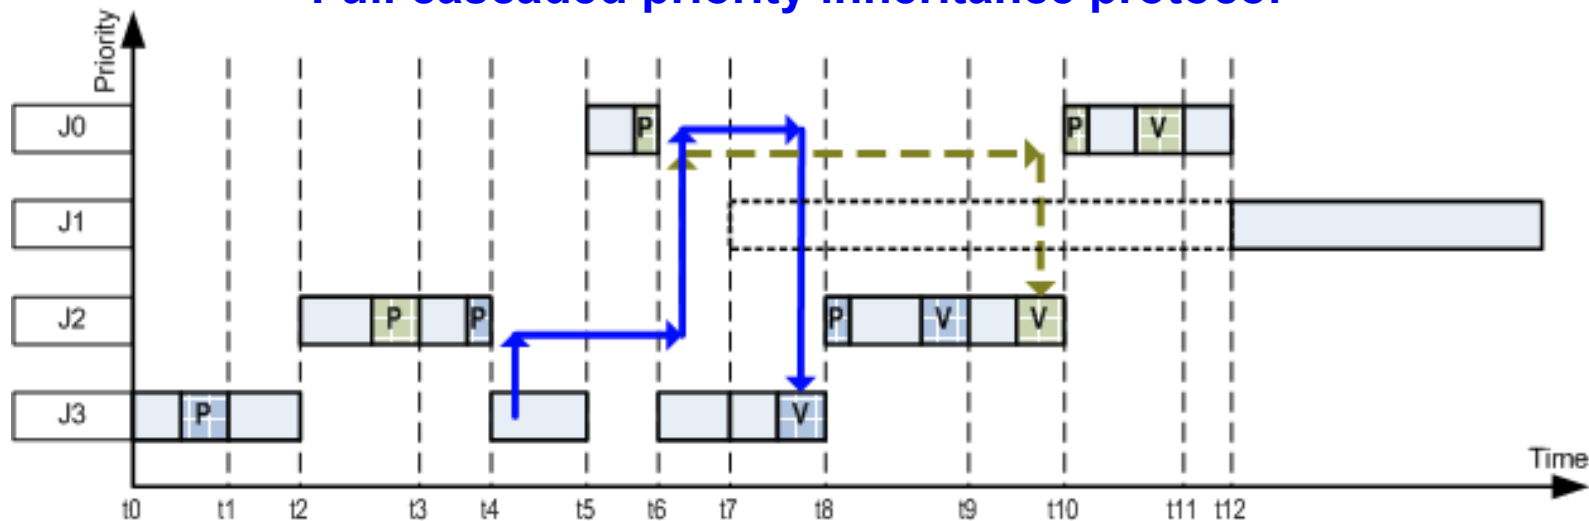


- Prevents priority inversion by temporarily raising the priority of the task holding the mutex.
- Reactive protocol. Priority is raised to the priority of the task that is waiting on the mutex, if that is greater.
- No support for multiple mutexes.

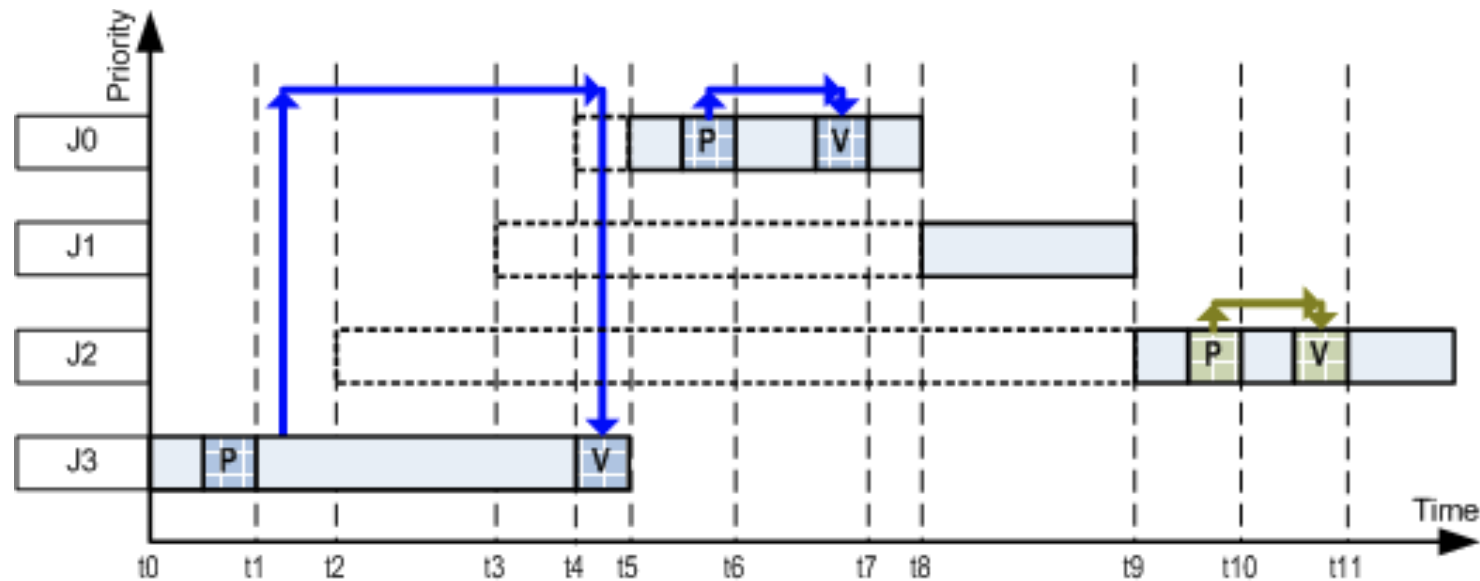
Priority Inheritance (with multiple mutexes)



Full cascaded priority inheritance protocol



Priority Ceiling Protocol

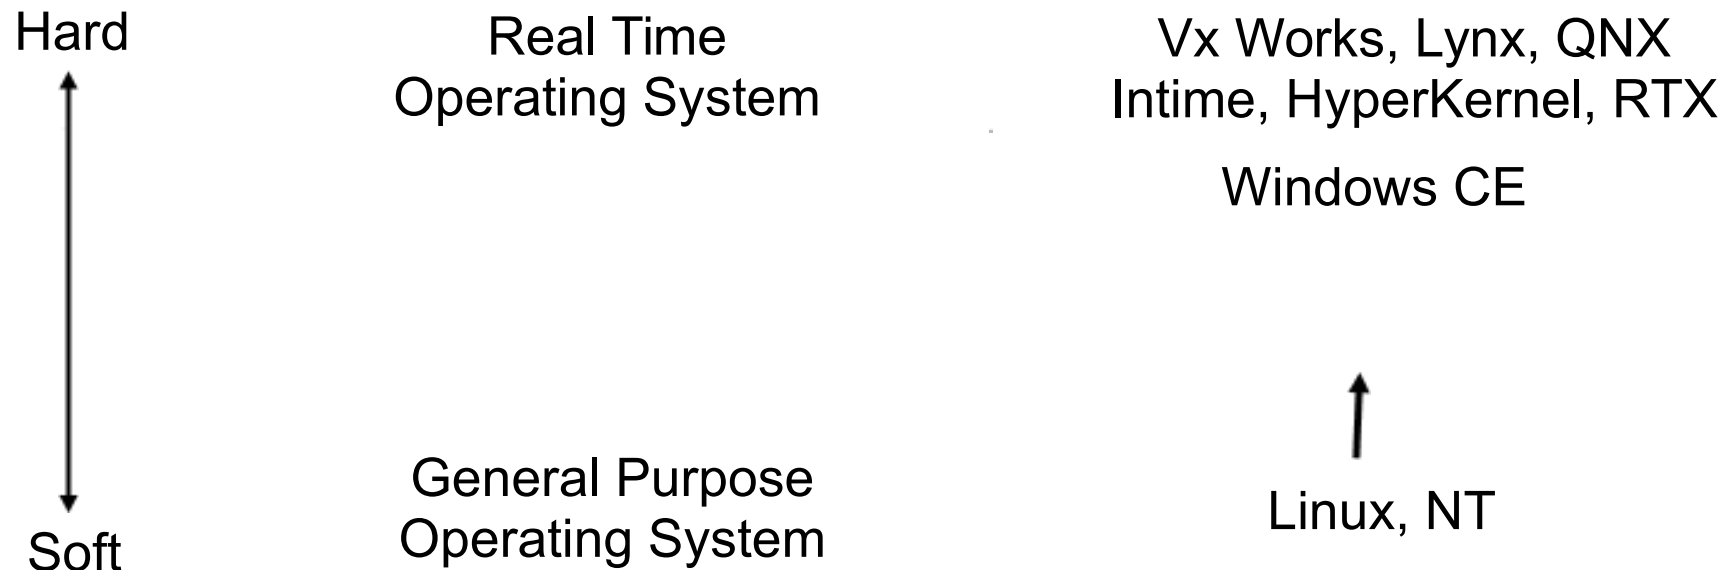


- Called **Priority Protect Protocol** in **POSIX**.
- Every mutex has a **priority ceiling**. The priority ceiling of a mutex should be greater than or equal to the priority of any task that might use it.
- **Proactive protocol**. When a task locks a mutex, its priority is immediately raised to the priority ceiling of the mutex.
- It can eliminate **deadlock**!

Overview of available RTOS's

- **Three categories of real-time operating systems:**

- Small, proprietary kernels. e.g. VRTX32, pSOS, VxWorks, Windows CE, MicroC-OS/III*
- Real-time Linux extensions: RT-Linux, Xenomai, RTAI
- Research kernels: MARS, ARTS, Spring, Polis, MicroC-OS/II

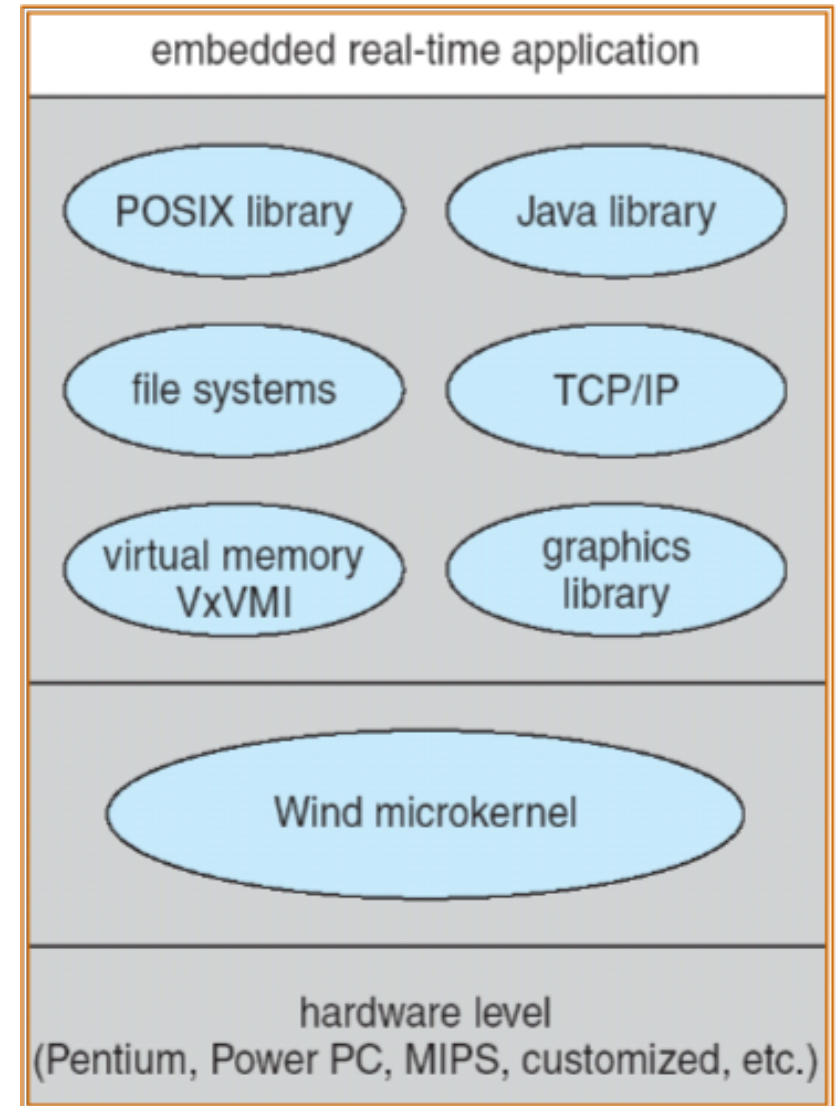


QNX Neutrino

- **Developed by Gordon Bell and the students at the University of Waterloo in 1980**
- **Supported on ARM, MIPS, Power PC, x86 and Pentium**
- **Micro-kernel based, and most of the OS is run in the form of a number of small tasks, known as servers**
- **Configurable to small size (64 K kernel ROM, 32 K kernel RAM)**
- **Supports Symmetric multiprocessing and strict priority-preemptive scheduling**
- **IEEE1003 real-time std compliant and POSIX threads**
- **Finds Applications in Embedded systems for over 20 years in mission and life critical systems, medical instrumentation, aviation and space systems, process control systems**
- **Avg Interrupt latency - 1.6us**

Vx Works

- **Developed by Wind River Systems of California**
- **Pentium, Motorola, Power PC, ARM**
- **Micro-kernel based**
- **Preemptive and non-preemptive scheduling**
- **Manages interrupts with bounded interrupt and dispatch latency times**
- **POSIX Compliant threads**
- **Shared memory and message passing for inter process communications**
- **Used in automobiles, routers, switches, Mars Pathfinder**
- **Avg interrupt latency - 1.7us**



LynxOS

- **Lynx is a Unix like real time operating system**
- **Developed for Motorola 68010, ported to x86, ARM, Power PC**
- **Support hard real time applications, due to extremely fast interrupt routines known as Multiple Priority Light Weight kernel Task-based Interrupt Handling**
- **Mostly used in embedded systems in avionics, aerospace, communications**

Microsoft Windows CE

- **Supported on x86, MIPS, ARM processors for embedded systems**
- **Supports threads, priority inheritance**
- **Non - POSIX compliant, 10% of Win32 APIs**
- **Applications can be developed on Visual Studio**
- **Windows Mobile, Pocket PC, Smart Phone are based on CE**
- **Avg interrupt latency - 2.4us**

Variants of Real-Time Linux Extensions

Real Time Linux (RTLinux)	Real Time Application Interface (RTAI)	Xenomai Framework
<ul style="list-style-type: none">– Developed at the New Mexico Institute of Mining and Technology– RTOS Micro kernel running entire Linux in fully preemptive mode– Runs special real-time tasks and interrupt handlers– FiFo, Shared memory, Semaphores. POSIX mutexes and threads– Avg interrupt latency - 15us– Used to control robots, data acquisition systems, manufacturing plants, and other time-sensitive instruments and machines	<ul style="list-style-type: none">– Developed by programmers at the Department of Aerospace Engineering, Milano– Adeos based patch over Linux kernel, with native real time tasks, interrupt handlers and services– Platforms - MIPS, x86-64, PowerPC, ARM– Semaphores, mailboxes, FIFOs, shared memory, RPCs– POSIX 1003.1c & POSIX 1003.1b(pqueues)– Avg interrupt response - 20us	<ul style="list-style-type: none">– Implementing and migrating real time applications– Based on emulators for proprietary RTOS interfaces/apis, such as VxWorks and pSOS– Linux-hosted dual kernel, with pure Adeos.– User space RT tasks– Platforms - x86, ARM, POWER, IA-64, Blackfin, nios